

The Forward-Forward Algorithm: Some Preliminary Investigations

G. Hinton

Presented by
Vishal V

Table

What's wrong with backprop?

Explaining FF

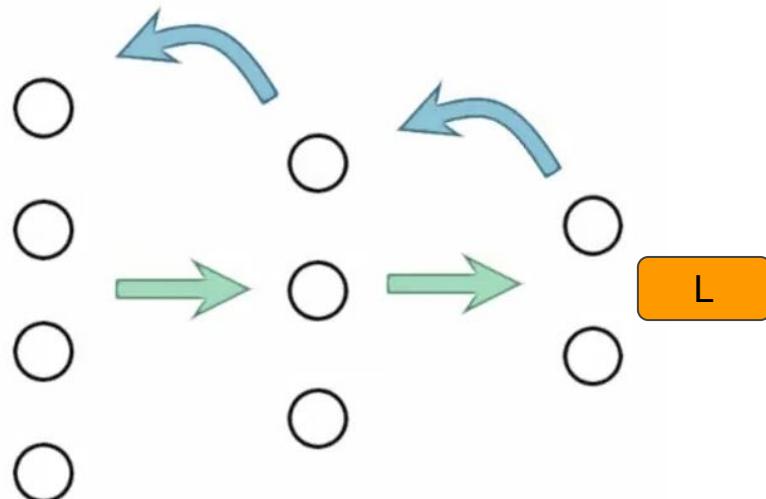
Approach and Results

Conclusion

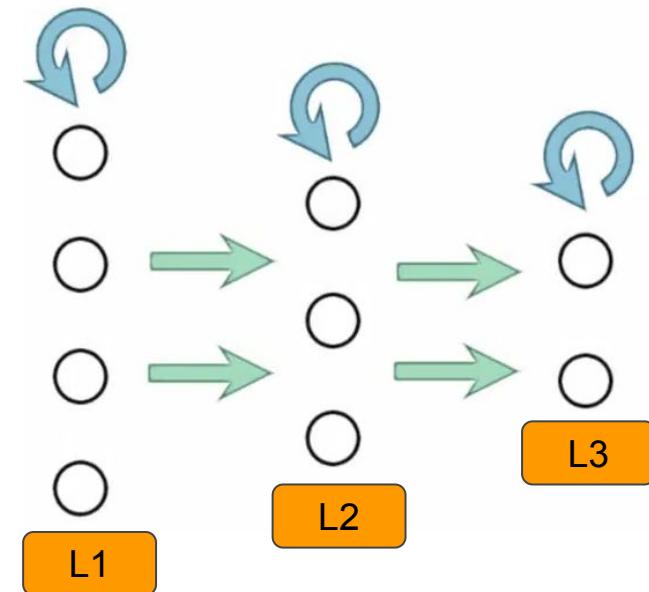
What's wrong with backprop

- Not Biologically Plausible
 - No evidence that brain store neural activities. For a backward pass like in backprop, storing activities would be necessary
 - Neurons do not wait for signals to pass through the entire brain to update connections. Brain needs to perform actions and learning in real-time without stopping
- Requires perfect knowledge of the forward pass
 - Weights and Biases can be updated only if we have total knowledge of forward and backward pass

Backpropagation



Forward-Forward



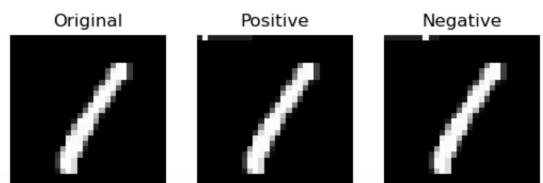
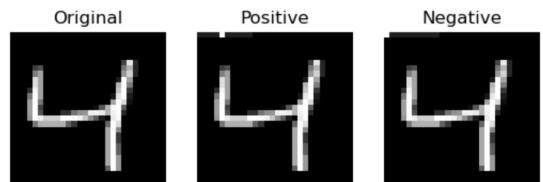
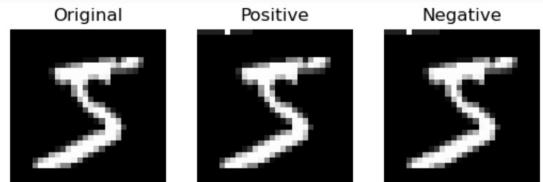
Loss function

$$\text{loss} = \frac{1}{N} \sum_{i=1}^N (\log(1 + \exp(-g_{\text{pos},i} + \text{threshold})) + \log(1 + \exp(g_{\text{neg},i} - \text{threshold})))$$

- N is the total number of elements in the tensors g_{pos} and g_{neg} .
- $g_{\text{pos},i}$ and $g_{\text{neg},i}$ are the i -th elements of the tensors g_{pos} and g_{neg} , respectively.
- threshold is a scalar value.
- The log and exp functions apply element-wise.
- The \sum notation indicates that we sum over all the elements resulting from the log operations, and then divide by N to compute the mean.

$$g = \frac{1}{M} \sum_{j=1}^M y_j^2$$

Positive and Negative data



Positive data

- Correctly labeled data
- Should increase the mean of squared activities above a threshold

For 0 original

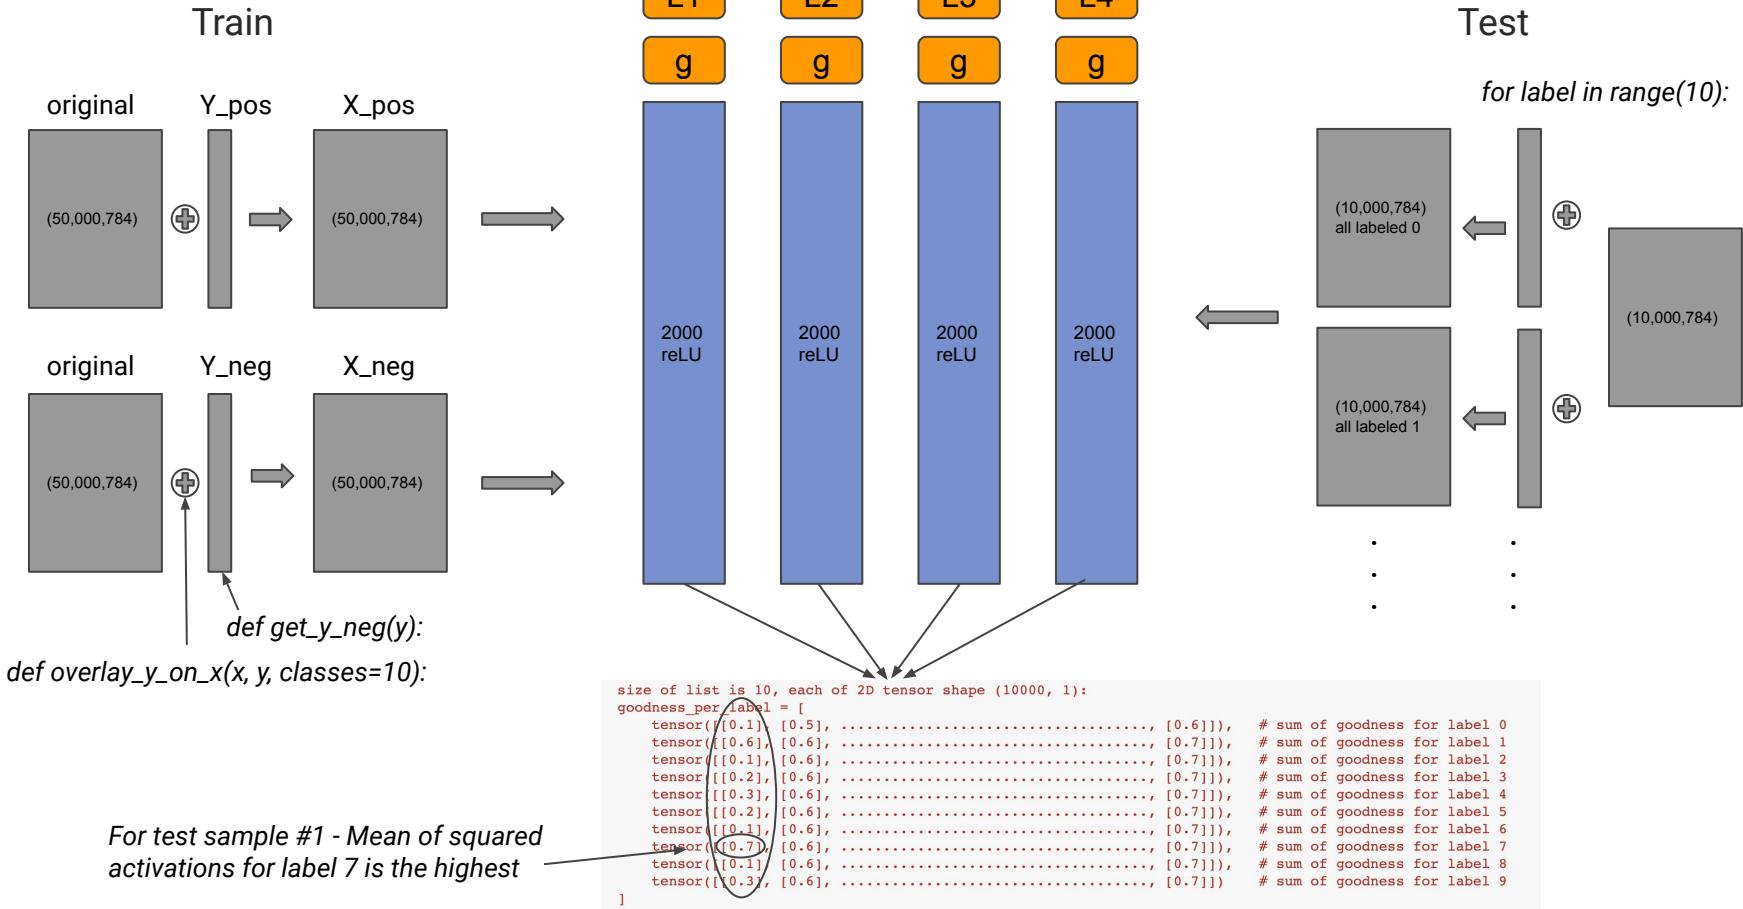


Negative data

- Incorrectly labeled data
- Should decrease the mean of squared activities above a threshold



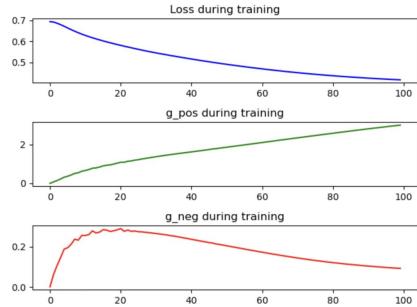
Fully Connected MNIST



Layerwise Loss and goodness during train

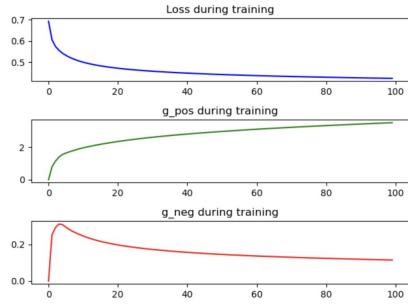
Threshold = 0

Layer 0



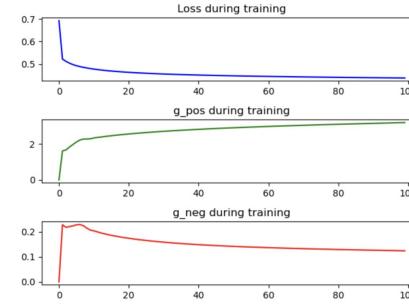
```
Loss at step 990: 0.4170325596469353
Loss at step 991: 0.41695263981819153
Loss at step 992: 0.41687288880348206
Loss at step 993: 0.41679325698906213
Loss at step 994: 0.4167138636112213
Loss at step 995: 0.41669526132794265
Loss at step 996: 0.416556132794265
Loss at step 997: 0.4164767861366272
Loss at step 998: 0.4163980782032013
```

Layer 1



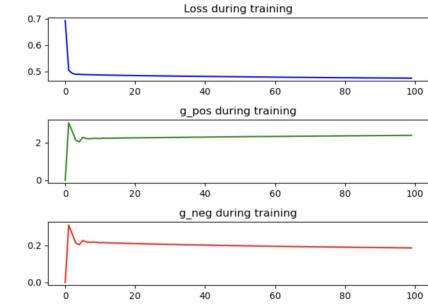
```
Loss at step 990: 0.42403945326805115
Loss at step 991: 0.42401471734046936
Loss at step 992: 0.423998181438876
Loss at step 993: 0.4239855298792
Loss at step 994: 0.4239405691623688
Loss at step 995: 0.423911592264175415
Loss at step 996: 0.4238913059234619
Loss at step 997: 0.4238667190749207
Loss at step 998: 0.4238421320915222
Loss at step 999: 0.42381757787787476
```

Layer 2



```
Loss at step 990: 0.4373338222503662
Loss at step 991: 0.43731990456581116
Loss at step 992: 0.43730592727661133
Loss at step 993: 0.437292595217098236
Loss at step 994: 0.437281217098236
Loss at step 995: 0.43726423382759094
Loss at step 996: 0.43725037574768066
Loss at step 997: 0.437236487865448
Loss at step 998: 0.4372226595878601
Loss at step 999: 0.4372088313102722
```

Layer 3



```
Loss at step 990: 0.4747256934642792
Loss at step 991: 0.47471755743026733
Loss at step 992: 0.4747093617916107
Loss at step 993: 0.4746983617916107
Loss at step 994: 0.47469303011894236
Loss at step 995: 0.4746849238872528
Loss at step 996: 0.47467678785324097
Loss at step 997: 0.47466862201690674
Loss at step 998: 0.4746604859828949
Loss at step 999: 0.47465240955352783
```

Method	Layers	Epochs	Train Accuracy	Test Accuracy
FF	(2000, 2000, 2000, 2000)	1000 epochs	93.29%	93.18%

FF vs backprop results

Method	Layers	Epochs	Train Accuracy	Test Accuracy	Threshold
FF	(2000, 2000, 2000, 2000)	1000 epochs	93.29%	93.18%	0
FF	(2000, 2000, 2000, 2000)	1000 epochs	93.46%	93.58%	2
FF	(2000, 2000, 2000, 2000)	1000 epochs	93.45%	93.58%	3
Backprop	(2000, 2000, 2000, 2000)	15 epochs	93.73%	94.36%	-

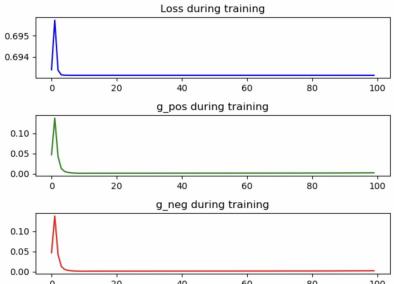
FF in convnets CIFAR-10

Network

```
# trial 7

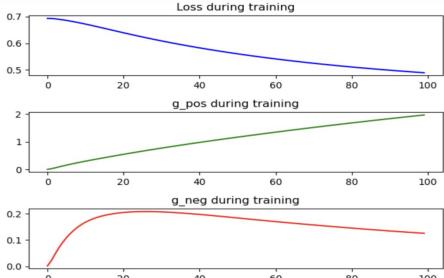
self.layers = nn.Sequential(
    Layer(3, 64, kernel_size=5, padding=2), # Layer 0 with 5x5 receptive field
    nn.ReLU(inplace=True),
    Layer(64, 128, kernel_size=5, padding=2, stride=2), # Layer 2 with 5x5 receptive field
    nn.ReLU(inplace=True),
    Layer(128, 256, kernel_size=5, padding=2, stride=2), # Layer 4 with 5x5 receptive field
    nn.ReLU(inplace=True),
    Layer(256, 10, kernel_size=1, stride=1), # Layer 6 with 1x1 receptive field
    nn.ReLU(inplace=True),
)
```

Layer 0



```
Loss at step 990: 0.6931461691856384
Loss at step 991: 0.6931461691856384
Loss at step 992: 0.6931461691856384
Loss at step 993: 0.6931461691856384
Loss at step 994: 0.6931461691856384
Loss at step 995: 0.6931461691856384
Loss at step 996: 0.6931461691856384
Loss at step 997: 0.6931461691856384
Loss at step 998: 0.6931461691856384
Loss at step 999: 0.6931461691856384
training layer: 1
training layer: 2
training layer: 3
training layer: 4
```

Layer 2

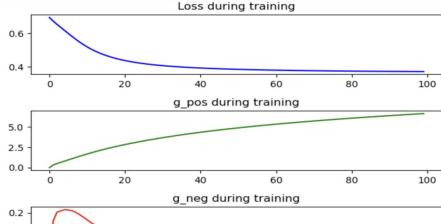


```
Loss at step 990: 0.48919624091947
Loss at step 991: 0.4890972375869751
Loss at step 992: 0.4890972375869751
Loss at step 993: 0.4889972375869751
Loss at step 994: 0.4889972375869751
Loss at step 995: 0.48880211168981
Loss at step 996: 0.48880211168981
Loss at step 997: 0.48880211168981
Loss at step 998: 0.48880211168981
Loss at step 999: 0.48880211168981
training layer: 5
training layer: 6
```

Forward Pass in a layer

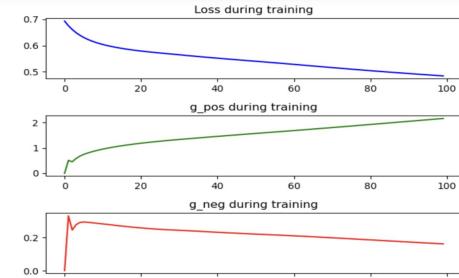
```
def forward(self, x):
    # L2 norm
    x_direction = x / (x.norm(2, 1, keepdim=True) + 1e-4)
    return self.relu(self.conv(x_direction))
```

Layer 4



```
Loss at step 990: 0.3722074627876282
Loss at step 991: 0.3722074627876282
Loss at step 992: 0.3722074627876282
Loss at step 993: 0.3722074627876282
Loss at step 994: 0.3722074627876282
Loss at step 995: 0.3722074627876282
Loss at step 996: 0.3722074627876282
Loss at step 997: 0.3722074627876282
Loss at step 998: 0.3722074627876282
Loss at step 999: 0.3722074627876282
training layer: 5
training layer: 6
```

Layer 6



```
Loss at step 990: 0.484397530557251
Loss at step 991: 0.48430374984712
Loss at step 992: 0.4842032194137573
Loss at step 993: 0.4841061532497406
Loss at step 994: 0.4840081333010104
Loss at step 995: 0.4839122593402862
Loss at step 996: 0.483815461397171
Loss at step 997: 0.4837186932563782
Loss at step 998: 0.4836186932563782
Loss at step 999: 0.48352542519569397
training layer: 7
```

Method	Layers	Epochs	Train Accuracy	Test Accuracy	Threshold
FF	trial #7	1000	9.90%	13.00%	0.002

Deeper convnets

trial 8

```

self.layers = nn.Sequential(
    Layer(3, 64, kernel_size=3, padding=1), # Layer 0
    nn.ReLU(inplace=True),
    Layer(64, 64, kernel_size=3, padding=1), # Layer 2
    nn.ReLU(inplace=True),

    nn.MaxPool2d(kernel_size=2, stride=2), # Layer 4

    Layer(64, 128, kernel_size=3, padding=1), # Layer 5
    nn.ReLU(inplace=True),
    Layer(128, 128, kernel_size=3, padding=1), # Layer 7
    nn.ReLU(inplace=True),

    nn.MaxPool2d(kernel_size=2, stride=2), # Layer 9

    Layer(128, 256, kernel_size=3, padding=1), # Layer 10
    nn.ReLU(inplace=True),
    Layer(256, 256, kernel_size=3, padding=1), # Layer 12
    nn.ReLU(inplace=True),
    Layer(256, 256, kernel_size=3, padding=1), # Layer 14
    nn.ReLU(inplace=True),

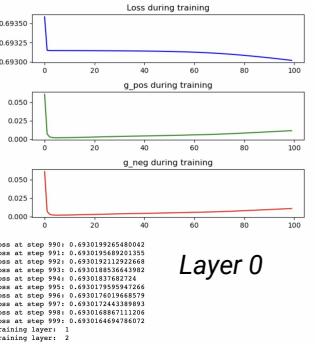
    nn.MaxPool2d(kernel_size=2, stride=2), # Layer 16

    Layer(256, 512, kernel_size=3, padding=1), # Layer 17
    nn.ReLU(inplace=True),
    Layer(512, 512, kernel_size=3, padding=1), # Layer 19
    nn.ReLU(inplace=True),
    Layer(512, 512, kernel_size=3, padding=1), # Layer 21
    nn.ReLU(inplace=True),

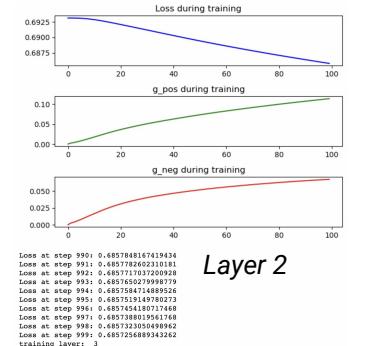
    nn.MaxPool2d(kernel_size=2, stride=2), # Layer 23

    Layer(512, 512, kernel_size=3, padding=1), # Layer 24
    nn.ReLU(inplace=True),
    Layer(512, 512, kernel_size=3, padding=1), # Layer 26
    nn.ReLU(inplace=True),
    Layer(512, 512, kernel_size=3, padding=1), # Layer 28
    nn.ReLU(inplace=True)
)
```

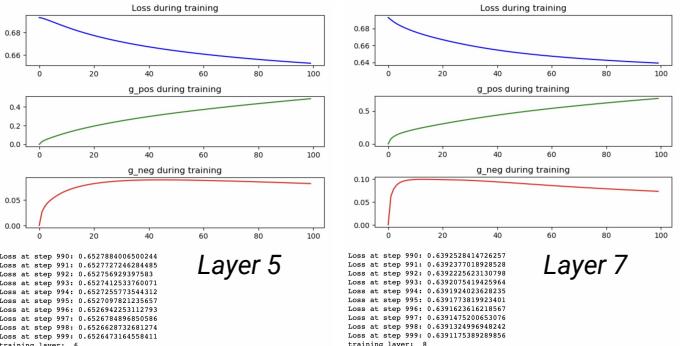
1



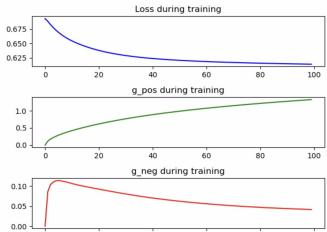
Layer 0



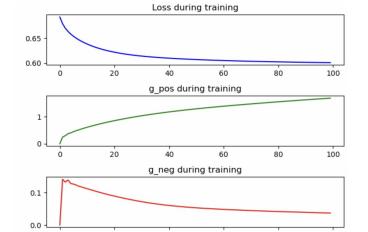
Layer 2



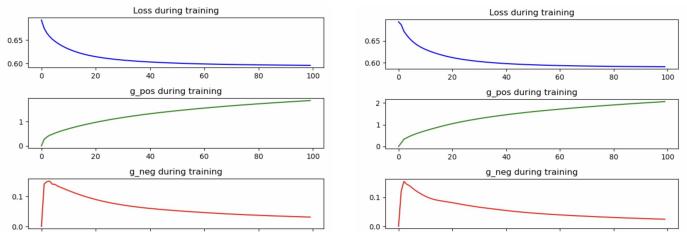
Layer 5



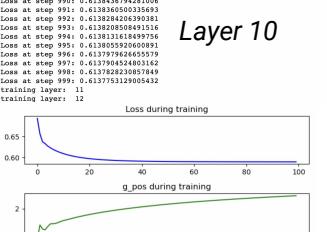
Layer 10



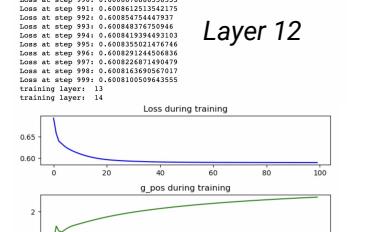
Layer 12



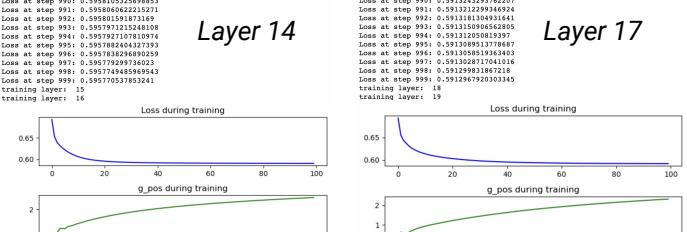
Layer 7



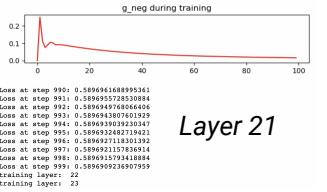
Layer 19



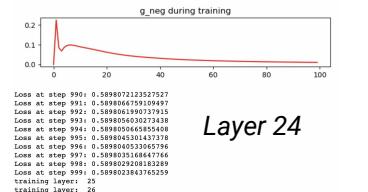
Layer 21



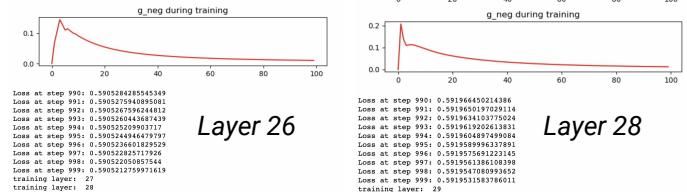
-aver 26



Layer 19



Layer 24



Layer 28

Wider convnets

```
# trial 9

self.layers = nn.Sequential(
    Layer(3, 64, kernel_size=5, padding=2), # Layer 0
    nn.ReLU(inplace=True),

    Layer(64, 128, kernel_size=5, padding=2, stride=2), # Layer 2
    nn.ReLU(inplace=True),

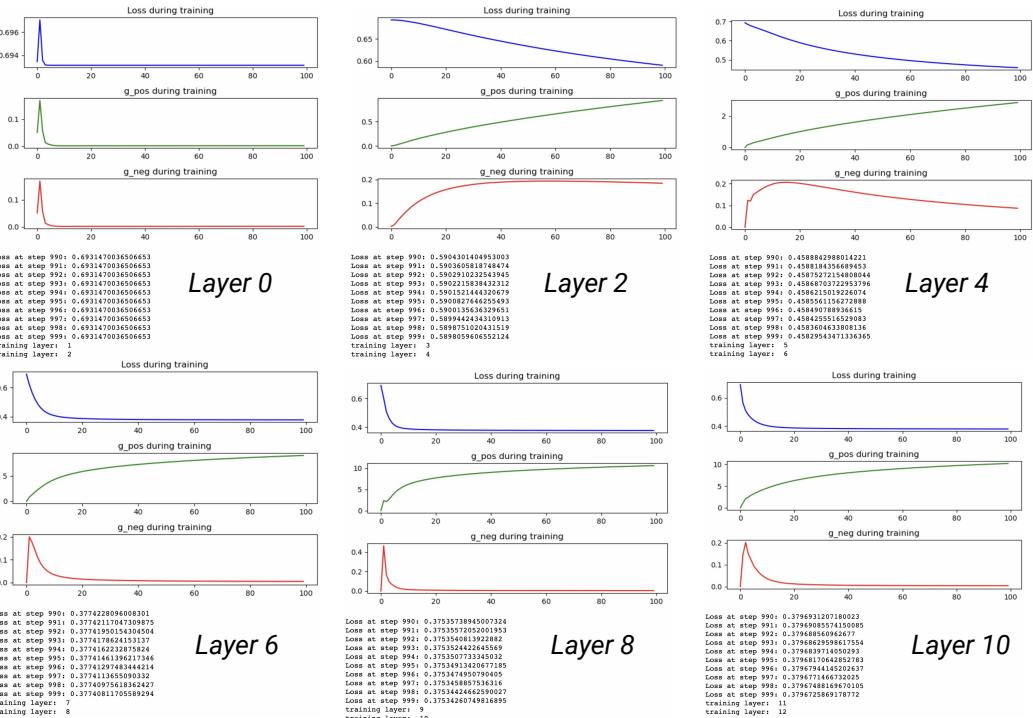
    Layer(128, 256, kernel_size=5, padding=2, stride=2), # Layer 4
    nn.ReLU(inplace=True),

    Layer(256, 512, kernel_size=5, padding=2, stride=2), # Layer 6
    nn.ReLU(inplace=True),

    Layer(512, 1024, kernel_size=5, padding=2, stride=2), # Layer 8
    nn.ReLU(inplace=True),

    Layer(1024, 2048, kernel_size=5, padding=2, stride=2), # Layer 10
    nn.ReLU(inplace=True),

    Layer(2048, 10, kernel_size=1, stride=1), # Layer 12
    nn.ReLU(inplace=True),
)
```

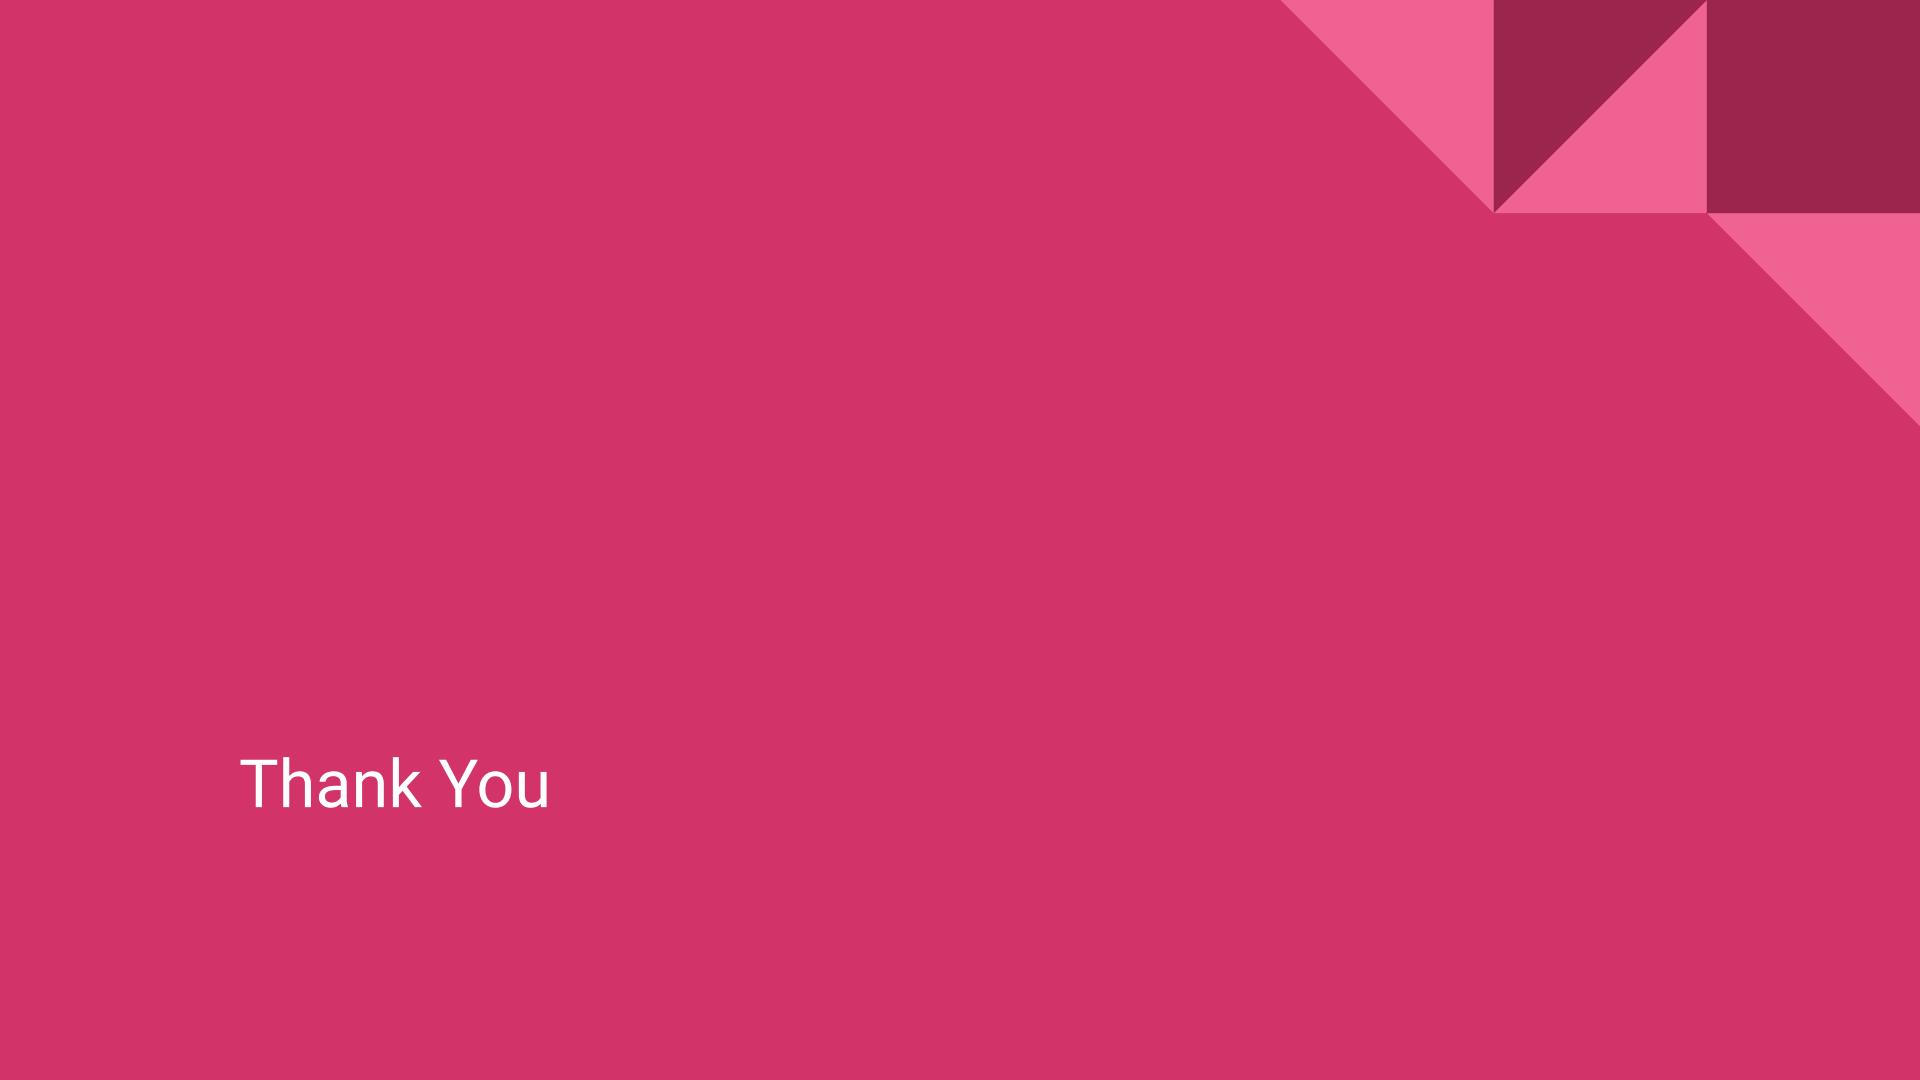


FF	trial #8	1000	9.90%	13.00%	0.002
FF	trial #9	1000	9.20%	10.00%	0.002

FF seems to be implausible for training convnets with the current loss and g function

Conclusion

- Explore other g function candidates - might solve for layerwise convnet training.
- Independent threshold params for each layer for better optimization.
- Try Biologically plausible activations - ReLU is used so far.



Thank You